


EgoLongQA question clusters (KMeans k=10, UMAP projection) — text-only vs multimodal encoder  
cluster agreement: adjusted Rand index = 0.16 (1 = identical partition, 0 = random)

sentence-transformers/all-MiniLM-L6-v2  
(text-only sentence encoder, silhouette=0.028)



google/siglip2-so400m-patch14-384  
(SigLIP-2 multimodal text tower, silhouette=0.005)

